

# CIFAR10 图片分类报告

# 内容提要

- 背景
- 模型和方法
- 实验分析与展示
- 总结












## 背景 - CIFAR10

- CIFAR10 是一个带标签的数据集
- 是一个图片分类任务，一共有 10 种不同的物品
  - 'plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck'
- 每一张图片大小为 $32 \times 32$ ，是彩色图片，3 通道(RGB)
- 训练集一共 50000 张图片，测试集 10000 张
- 数据集及网站

<http://www.cs.toronto.edu/~kriz/cifar.html>

# 背景 - CIFAR10

## · CIFAR10 现在的成绩

Result	Method	Venue	Details
96.53%	<a href="#">Fractional Max-Pooling</a> 	arXiv 2015	<a href="#">Details</a>
95.59%	<a href="#">Striving for Simplicity: The All Convolutional Net</a> 	ICLR 2015	<a href="#">Details</a>
94.16%	<a href="#">All you need is a good init</a> 	ICLR 2016	<a href="#">Details</a>
94%	<a href="#">Lessons learned from manually classifying CIFAR-10</a> 	unpublished 2011	<a href="#">Details</a>
93.95%	<a href="#">Generalizing Pooling Functions in Convolutional Neural Networks: Mixed, Gated, and Tree</a> 	AISTATS 2016	<a href="#">Details</a>
93.72%	<a href="#">Spatially-sparse convolutional neural networks</a> 	arXiv 2014	
93.63%	<a href="#">Scalable Bayesian Optimization Using Deep Neural Networks</a> 	ICML 2015	
93.57%	<a href="#">Deep Residual Learning for Image Recognition</a> 	arXiv 2015	<a href="#">Details</a>
93.45%	<a href="#">Fast and Accurate Deep Network Learning by Exponential Linear Units</a> 	arXiv 2015	<a href="#">Details</a>
93.34%	<a href="#">Universum Prescription: Regularization using Unlabeled Data</a> 	arXiv 2015	
93.25%	<a href="#">Batch-normalized Maxout Network in Network</a> 	arXiv 2015	<a href="#">Details</a>

[http://rodrigob.github.io/are\\_we\\_there\\_yet/build/classification\\_datasets\\_results.html#43494641522d3130](http://rodrigob.github.io/are_we_there_yet/build/classification_datasets_results.html#43494641522d3130)

- 背景


- 模型和方法

- 实验分析与展示

- 总结

# 模型和方法 - 模型 - VGG16

- VGG - Visual Geometry Group(University of Oxford)
- VGG16
  - 13卷积层 + 3全连接层



# 模型和方法 - 模型 - VGG16

- 初始模型
  - 数据初始化
    - 归一化 [-1,1]
  - 每一个神经网络节点
    - 归一化 [-1,1]
    - 激活函数 Relu
- 损失函数：交叉熵
- 优化器：梯度下降法
  - 学习率：0.001

# 模型和方法 - 模型 - VGG16

- 初始模型
  - 卷积层  $C(\text{inChannels}, \text{outChannels})$ 
    - $\text{kernelSize} = 3 \times 3$ ,  $\text{stride} = 1$ ,  $\text{padding} = 1$
  - 池化层  $M$ 
    - $\text{maxpool}$ ,  $\text{size} = 2 \times 2$ ,  $\text{stride} = 2$
  - 全连接层  $\text{FC}(\text{inputSize}, \text{outputSize})$ 
    - 由于图片较小全连接层只保留一层



# 模型和方法 - 模型 - VGG16


· 输入:  $32*32*3$

# 1.  $3*32*32$  -> 卷积核个数:64  
# 2.  $64*32*32$  -> 卷积核个数:64  
# 3.  $64*32*32$  -> maxpool: $2*2$   
# 4.  $64*16*16$  -> 卷积核个数:128  
# 5.  $128*16*16$  -> 卷积核个数:128  
# 6.  $128*16*16$  -> maxpool: $2*2$   
# 7.  $128*8*8$  -> 卷积核个数:256  
# 8.  $256*8*8$  -> 卷积核个数:256  
# 9.  $256*8*8$  -> 卷积核个数:256  
#10.  $256*8*8$  -> maxpool: $2*2$


#11.  $256*4*4$  -> 卷积核个数:512  
#12.  $512*4*4$  -> 卷积核个数:512  
#13.  $512*4*4$  -> 卷积核个数:512  
#14.  $512*4*4$  -> maxpool: $2*2$   
#15.  $512*2*2$  -> 卷积核个数:512  
#16.  $512*2*2$  -> 卷积核个数:512  
#17.  $512*2*2$  -> 卷积核个数:512  
#18.  $512*2*2$  -> maxpool: $2*2$   
#19.  $512*1*1$  -> Linear:512->10

# 模型和方法 - 模型 - VGG16

· 输入:  $32*32*3$



# 模型和方法 - 模型 - VGG16



# 模型和方法 - 方法 - pytorch

## · 数据初始化：读入 + 正则化

```
# Normalize 正则化(预处理)
# torchvision.transforms.Normalize(mean, std, inplace=False)
# input[channel] = (input[channel] - mean[channel]) / std[channel]
# 如下示例,输入原本为 [0,1] -> [-1,1]
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

# 训练集
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

# 测试集
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)
```

## · 注意每一组有 4 个数据

# 模型和方法 - 方法 - pytorch

## · 卷积层 + 池化层

```
... # 网络层
... self.layers = nn.Sequential(
...     # CCP
...     nn.Conv2d(3, 64, 3, stride=1, padding=1),
...     # 正则化
...     nn.BatchNorm2d(64),
...     # 激活函数
...     nn.ReLU(inplace=True),
...     nn.Conv2d(64, 64, 3, stride=1, padding=1),
...     nn.BatchNorm2d(64),
...     nn.ReLU(inplace=True),
...     nn.MaxPool2d(kernel_size=2, stride=2),
...     # CCP
...     nn.Conv2d(64, 128, 3, stride=1, padding=1),
...     nn.BatchNorm2d(128),
...     nn.ReLU(inplace=True),
...     nn.Conv2d(128, 128, 3, stride=1, padding=1),
...     nn.BatchNorm2d(128),
...     nn.ReLU(inplace=True),
...     nn.MaxPool2d(kernel_size=2, stride=2),
...     # CCCP
...     nn.Conv2d(128, 256, 3, stride=1, padding=1),
...     nn.BatchNorm2d(256),
...     nn.ReLU(inplace=True),
...     nn.Conv2d(256, 256, 3, stride=1, padding=1),
...     nn.BatchNorm2d(256),
...     nn.ReLU(inplace=True),
...     nn.Conv2d(256, 256, 3, stride=1, padding=1),
...     nn.BatchNorm2d(256),
...     nn.ReLU(inplace=True),
...     nn.MaxPool2d(kernel_size=2, stride=2),
```

```
...     # CCCP
...     nn.Conv2d(256, 512, 3, stride=1, padding=1),
...     nn.BatchNorm2d(512),
...     nn.ReLU(inplace=True),
...     nn.Conv2d(512, 512, 3, stride=1, padding=1),
...     nn.BatchNorm2d(512),
...     nn.ReLU(inplace=True),
...     nn.Conv2d(512, 512, 3, stride=1, padding=1),
...     nn.BatchNorm2d(512),
...     nn.ReLU(inplace=True),
...     nn.MaxPool2d(kernel_size=2, stride=2),
...     # CCCP
...     nn.Conv2d(512, 512, 3, stride=1, padding=1),
...     nn.BatchNorm2d(512),
...     nn.ReLU(inplace=True),
...     nn.Conv2d(512, 512, 3, stride=1, padding=1),
...     nn.BatchNorm2d(512),
...     nn.ReLU(inplace=True),
...     nn.Conv2d(512, 512, 3, stride=1, padding=1),
...     nn.BatchNorm2d(512),
...     nn.ReLU(inplace=True),
...     nn.MaxPool2d(kernel_size=2, stride=2)
... )
```

# 模型和方法 - 方法 - pytorch

## · 全连接层

```
···# 网络层的构建
···def forward(self, x):
·····output = self.layers(x)
·····# 一个 batch 有4个样本
·····output = output.view(4, 512)
·····output = self.classifier(output)
·····return output
```

# 模型和方法 - 方法 - pytorch

- 数据初始化: argmax

```
def calcAcc(isTrain):  
    correct = 0  
    total = 0  
    with torch.no_grad():  
        for data in trainloader if isTrain else testloader:  
            images, labels = data  
            if useGPU:  
                images = images.cuda()  
                labels = labels.cuda()  
            outputs = net(images)  
            _, predicted = torch.max(outputs.data, 1)  
            total += labels.size(0)  
            correct += (predicted == labels).sum().item()  
    return 1.0 * correct / total
```

# 模型和方法 - 方法 - pytorch

- 其他

- GPU 加速

```
# GPU
useGPU = torch.cuda.is_available()
if useGPU:
    net = net.cuda()
    net = torch.nn.DataParallel(net, device_ids=range(torch.cuda.device_count()))
    torch.backends.cudnn.benchmark = True
    criterion.cuda()
    # optimizer.cuda()
    print("USING GPU(CUDA)!")
else:
    print("USING CPU!")
```




# 模型和方法 - 方法 - pytorch

- 之后根据过拟合情况使用了其他优化方案
  - 在之后版本迭代的过程中给出
  - version1(初始版本) -> version6

- 背景
- 模型和方法
- 实验分析与展示
- 总结

# 实验分析与展示 - version1

· epoch [0-12]



·过拟合严重

·测试集最高准确率

·82.44%

# 实验分析与展示 - version2

- FractionalMaxPool

- 将 MaxPool 换成 FractionalMaxPool

- Benjamin Graham “Fractional Max-Pooling”


- <https://arxiv.org/pdf/1412.6071.pdf>

- 分数池化 -> 提高泛化能力，减少过拟合

```
... nn.Conv2d(256, 256, 3, stride=1, padding=1),
... nn.BatchNorm2d(256),
... nn.ReLU(inplace=True),
... nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
... # CCCP
... nn.Conv2d(256, 512, 3, stride=1, padding=1),
... nn.BatchNorm2d(512),
... nn.ReLU(inplace=True),
... nn.Conv2d(512, 512, 3, stride=1, padding=1),
... nn.BatchNorm2d(512),
... nn.ReLU(inplace=True),
... nn.Conv2d(512, 512, 3, stride=1, padding=1),
... nn.BatchNorm2d(512),
... nn.ReLU(inplace=True),
... nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
...
```

# 实验分析与展示 - version2

· epoch [0-24]



·测试集最高准确率

·83.368%

·过拟合仍然严重


# 实验分析与展示 - version3(FractionalMaxPool + dropout)

- FractionalMaxPool + dropout
- 加入 dropout -> 减少过拟合
  - 在某些层上随机删除一些边
- 在最后两次最大池化后加入 20% dropout

```
.....# CCCP
.....# dropout-v3
.....nn.Dropout(p=0.2),
.....# dropout-v3
.....nn.Conv2d(512,512,3, stride=1, padding=1),
.....nn.BatchNorm2d(512),
.....nn.ReLU(inplace=True),
.....nn.Conv2d(512,512,3, stride=1, padding=1),
.....nn.BatchNorm2d(512),
.....nn.ReLU(inplace=True),
.....nn.Conv2d(512,512,3, stride=1, padding=1),
.....nn.BatchNorm2d(512),
.....nn.ReLU(inplace=True),
.....nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
.....# dropout-v3
.....nn.Dropout(p=0.2)
.....# dropout-v3
```

# 实验分析与展示 - version3

· epoch [0-19]



·测试集最高准确率

·83.309%

·效果并不好

# 实验分析与展示 - version4

- FractionalMaxPool + dropout + enhanced pre-process
- 加强对图像的预处理
  - 50% 的概率将图片进行竖直翻转
  - 将图片进行随机裁剪

```
transform1 = transforms.Compose(  
    ... # 随机裁剪 + padding  
    ... transforms.RandomCrop(32, padding=4),  
    ... # 50% 可能性竖直翻转  
    ... transforms.RandomHorizontalFlip(),  
    ... [transforms.ToTensor(),  
    ... transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])  
  
transform2 = transforms.Compose(  
    ... [transforms.ToTensor(),  
    ... transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```




## 实验分析与展示 - version4

- 随即裁剪 + padding 之后输出变成  $40*40*3$
- 不影响通道数
  - 随机裁剪不影响通道数
  - 池化不影响通道数
- 但是不影响主体结构，仍然通过 5 次池化变成  $1*1$ 
  - 由于始终是方形，只写一维
  - $40 \rightarrow 20 \rightarrow 10 \rightarrow 5 \rightarrow 2 \rightarrow 1$

# 实验分析与展示 - version4

· epoch [0-115]



·测试集最高准确率

·90.32%

·效果拔群

·准确率稳定在 90%

# 实验分析与展示 - version5

- FractionalMaxPool + dropout + enhanced pre-process  
+ 非常规正则化


```
transform1 = transforms.Compose([
    ... # 随机裁剪 + padding
    ... transforms.RandomCrop(32, padding=4),
    ... # 50% 可能性竖直翻转
    ... transforms.RandomHorizontalFlip(),
    ... transforms.ToTensor(),
    ... transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)))

transform2 = transforms.Compose(
    ... [transforms.ToTensor(),
    ... transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))])
```

<https://github.com/kuangliu/pytorch-cifar>

# 实验分析与展示 - version5

· epoch [0-50]



·测试集最高准确率

·88.85%

·训练集正确率还不高

97%

可以继续训练

# 实验分析与展示 - version6

- FractionalMaxPool + enhanced pre-process + 非常规正则化 + 强化版dropout + 最后再加一层全连接层

<http://torch.ch/blog/2015/07/30/cifar.html>

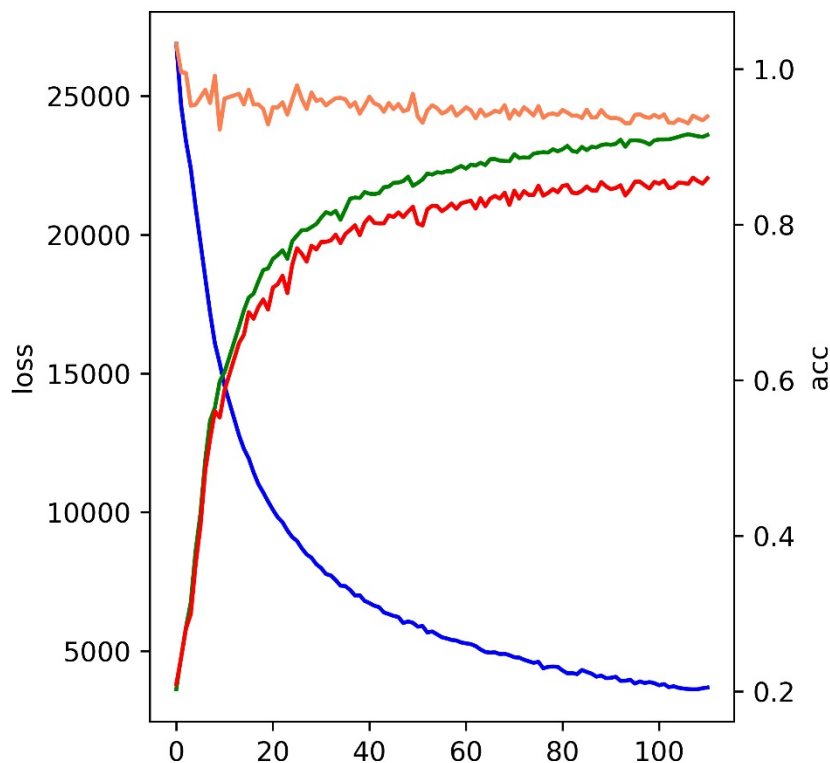
```

...# CCP
...nn.Conv2d(3, 64, 3, stride=1, padding=1),
...# 正则化
...nn.BatchNorm2d(64),
...# 激活函数
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.3), # v6
...nn.Conv2d(64, 64, 3, stride=1, padding=1),
...nn.BatchNorm2d(64),
...nn.ReLU(inplace=True),
...nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
...# CCP
...nn.Conv2d(64, 128, 3, stride=1, padding=1),
...nn.BatchNorm2d(128),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.Conv2d(128, 128, 3, stride=1, padding=1),
...nn.BatchNorm2d(128),
...nn.ReLU(inplace=True),
...nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
...# CCCP
...nn.Conv2d(128, 256, 3, stride=1, padding=1),
...nn.BatchNorm2d(256),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.Conv2d(256, 256, 3, stride=1, padding=1),
...nn.BatchNorm2d(256),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.Conv2d(256, 256, 3, stride=1, padding=1),
...nn.BatchNorm2d(256),
...nn.ReLU(inplace=True),
...nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
...nn.BatchNorm2d(512),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.Conv2d(512, 512, 3, stride=1, padding=1),
...nn.BatchNorm2d(512),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.Conv2d(512, 512, 3, stride=1, padding=1),
...nn.BatchNorm2d(512),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
...# CCCP
...nn.Conv2d(512, 512, 3, stride=1, padding=1),
...nn.BatchNorm2d(512),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.Conv2d(512, 512, 3, stride=1, padding=1),
...nn.BatchNorm2d(512),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.Conv2d(512, 512, 3, stride=1, padding=1),
...nn.BatchNorm2d(512),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...nn.FractionalMaxPool2d(kernel_size=2, output_ratio=(0.5, 0.5)),
...nn.BatchNorm2d(512),
...nn.ReLU(inplace=True),
...nn.Dropout(p=0.4), # v6
...# v6 FC
...self.fc1 = nn.Linear(512, 512)
...self.fc2 = nn.BatchNorm2d(512)
...self.fc3 = nn.ReLU()
...self.fc4 = nn.Dropout(p=0.5) # v6
...# Linear
...self.classifier = nn.Linear(512, 10)

```

# 实验分析与展示 - version6

· epoch [0-110]



<http://torch.ch/blog/2015/07/30/cifar.html>

·测试集最高准确率

·85.98%

·训练速度变慢

·参数多

·训练集正确率 91%

·还能继续训练

·GPU太贵 + 时间太长

·资料显示正确率可达

·92.45%

- 背景
- 模型和方法
- 实验分析与展示
- 总结

# 总结

- VGG16
- 尝试了不同的减少过拟合的方案
  - FractionalMaxPool
  - dropout
  - 增强数据预处理
  - 非标准正则化



# 总结

- 最终提交版本为version4
  - FractionalMaxPool + dropout + enhanced pre-process
- 不选 version5,version6
  - GPU太贵以及训练时间过长未完成训练

·最终提交版本为version4,epoch=107

[107]

loss: 381.36252677440643

train acc: 0.99226

test acc: 0.9032